

**NAME**

bridge – show / manipulate bridge addresses and devices

**SYNOPSIS**

**bridge** [ *OPTIONS* ] *OBJECT* { *COMMAND* | **help** }

*OBJECT* := { **link** | **fdb** | **mdb** | **vlan** | **monitor** }

*OPTIONS* := { **-V**[*ersion*] | **-s**[*tatistics*] | **-n**[*etns*] *name* | **-b**[*atch*] *filename* | **-c**[*lor*] | **-p**[*retty*] | **-j**[*son*] | **-o**[*neline*] }

**bridge link set dev** *DEV* [ **cost** *COST* ] [ **priority** *PRIO* ] [ **state** *STATE* ] [ **guard** { **on** | **off** } ] [ **hairpin** { **on** | **off** } ] [ **fastleave** { **on** | **off** } ] [ **root\_block** { **on** | **off** } ] [ **learning** { **on** | **off** } ] [ **learning\_sync** { **on** | **off** } ] [ **flood** { **on** | **off** } ] [ **hwmode** { **vepa** | **veb** } ] [ **mcast\_flood** { **on** | **off** } ] [ **mcast\_to\_unicast** { **on** | **off** } ] [ **neigh\_suppress** { **on** | **off** } ] [ **vlan\_tunnel** { **on** | **off** } ] [ **isolated** { **on** | **off** } ] [ **backup\_port** *DEVICE* ] [ **nobackup\_port** ] [ **self** ] [ **master** ]

**bridge link** [ **show** ] [ **dev** *DEV* ]

**bridge fdb** { **add** | **append** | **del** | **replace** } *LLADDR* **dev** *DEV* { **local** | **static** | **dynamic** } [ **self** ] [ **master** ] [ **router** ] [ **use** ] [ **extern\_learn** ] [ **sticky** ] [ **dst** *IPADDR* ] [ **src\_vni** *SRC\_VNI* ] [ **vni** *VNI* ] [ **port** *PORT* ] [ **via** *DEVICE* ]

**bridge fdb** [ **show** ] [ **dev** *DEV* ] [ **br** *BRDEV* ] [ **brport** *DEV* ] [ **vlan** *VID* ] [ **state** *STATE* ]

**bridge mdb** { **add** | **del** } **dev** *DEV* **port** *PORT* **grp** *GROUP* [ **permanent** | **temp** ] [ **vid** *VID* ]

**bridge mdb show** [ **dev** *DEV* ]

**bridge vlan** { **add** | **del** } **dev** *DEV* **vid** *VID* [ **tunnel\_info** *TUNNEL\_ID* ] [ **pvid** ] [ **untagged** ] [ **self** ] [ **master** ]

**bridge vlan** [ **show** | **tunnelshow** ] [ **dev** *DEV* ]

**bridge monitor** [ **all** | **neigh** | **link** | **mdb** ]

**OPTIONS**

**-V, -Version**

print the version of the **bridge** utility and exit.

**-s, -stats, -statistics**

output more information. If this option is given multiple times, the amount of information increases. As a rule, the information is statistics or some time values.

**-d, -details**

print detailed information about MDB router ports.

**-n, -net, -netns** <NETNS>

switches **bridge** to the specified network namespace *NETNS*. Actually it just simplifies executing of:

**ip netns exec** *NETNS* **bridge** [ *OPTIONS* ] *OBJECT* { *COMMAND* | **help** }

to

**bridge** -n[etns] *NETNS* [ *OPTIONS* ] *OBJECT* { *COMMAND* | **help** }

**-b, -batch** <FILENAME>

Read commands from provided file or standard input and invoke them. First failure will cause termination of bridge command.

**-force** Don't terminate bridge command on errors in batch mode. If there were any errors during execution of the commands, the application return code will be non zero.

**-c[*color*]** [= { **always** | **auto** | **never** }

Configure color output. If parameter is omitted or **always**, color output is enabled regardless of stdout state. If parameter is **auto**, stdout is checked to be a terminal before enabling color output. If parameter is **never**, color output is disabled. If specified multiple times, the last one takes precedence. This flag is ignored if **-json** is also given.

**-j, -json**

Output results in JavaScript Object Notation (JSON).

**-p, -pretty**

When combined with -j generate a pretty JSON output.

**-o, -oneline**

output each record on a single line, replacing line feeds with the '\n' character. This is convenient when you want to count records with **wc**(1) or to **grep**(1) the output.

## BRIDGE - COMMAND SYNTAX

### *OBJECT*

**link** - Bridge port.

**fdb** - Forwarding Database entry.

**mdb** - Multicast group database entry.

**vlan** - VLAN filter list.

### *COMMAND*

Specifies the action to perform on the object. The set of possible actions depends on the object type. As a rule, it is possible to **add**, **delete** and **show** (or **list**) objects, but some objects do not allow all of these operations or have some additional commands. The **help** command is available for all objects. It prints out a list of available commands and argument syntax conventions.

If no command is given, some default command is assumed. Usually it is **list** or, if the objects of this class cannot be listed, **help**.

**bridge link - bridge port**

**link** objects correspond to the port devices of the bridge.

The corresponding commands set and display port status and bridge specific attributes.

**bridge link set - set bridge specific attributes on a port**

**dev** *NAME*

interface name of the bridge port

**cost** *COST*

the STP path cost of the specified port.

**priority** *PRIOR*

the STP port priority. The priority value is an unsigned 8-bit quantity (number between 0 and 255). This metric is used in the designated port and root port selection algorithms.

**state** *STATE*

the operation state of the port. This is primarily used by user space STP/RSTP implementation. One may enter a lowercased port state name, or one of the numbers below. Negative inputs are ignored, and unrecognized names return an error.

**0** - port is DISABLED. Make this port completely inactive.

**1** - STP LISTENING state. Only valid if STP is enabled on the bridge. In this state the port listens for STP BPDUs and drops all other traffic frames.

**2** - STP LEARNING state. Only valid if STP is enabled on the bridge. In this state the port will accept traffic only for the purpose of updating MAC address tables.

**3** - STP FORWARDING state. Port is fully active.

**4** - STP BLOCKING state. Only valid if STP is enabled on the bridge. This state is used during the STP election process. In this state, port will only process STP BPDUs.

**guard on or guard off**

Controls whether STP BPDUs will be processed by the bridge port. By default, the flag is turned off allowing BPDUs to be processed. Turning this flag on will cause the port to stop processing STP BPDUs.

**hairpin on or hairpin off**

Controls whether traffic may be sent back out of the port on which it was received. By default, this flag is turned off and the bridge will not forward traffic back out of the receiving port.

**fastleave on or fastleave off**

This flag allows the bridge to immediately stop multicast traffic on a port that receives IGMP Leave message. It is only used with IGMP snooping enabled on the bridge. By default the flag is

off.

**root\_block on or root\_block off**

Controls whether a given port is allowed to become root port or not. Only used when STP is enabled on the bridge. By default the flag is off.

**learning on or learning off**

Controls whether a given port will learn MAC addresses from received traffic or not. If learning is off, the bridge will end up flooding any traffic for which it has no FDB entry. By default this flag is on.

**learning\_sync on or learning\_sync off**

Controls whether a given port will sync MAC addresses learned on device port to bridge FDB.

**flood on or flood off**

Controls whether a given port will flood unicast traffic for which there is no FDB entry. By default this flag is on.

**hwmode**

Some network interface cards support HW bridge functionality and they may be configured in different modes. Currently supported modes are:

**vepa** - Data sent between HW ports is sent on the wire to the external switch.

**veb** - bridging happens in hardware.

**mcast\_flood on or mcast\_flood off**

Controls whether a given port will flood multicast traffic for which there is no MDB entry. By default this flag is on.

**mcast\_to\_unicast on or mcast\_to\_unicast off**

Controls whether a given port will replicate packets using unicast instead of multicast. By default this flag is off.

**neigh\_suppress on or neigh\_suppress off**

Controls whether neigh discovery (arp and nd) proxy and suppression is enabled on the port. By default this flag is off.

**vlan\_tunnel on or vlan\_tunnel off**

Controls whether vlan to tunnel mapping is enabled on the port. By default this flag is off.

**isolated on or isolated off**

Controls whether a given port will be isolated, which means it will be able to communicate with non-isolated ports only. By default this flag is off.

**backup\_port DEVICE**

If the port loses carrier all traffic will be redirected to the configured backup port

**nobackup\_port**

Removes the currently configured backup port

**self** link setting is configured on specified physical device

**master** link setting is configured on the software bridge (default)

**-t, -timestamp**

display current time when using monitor option.

**bridge link show - list bridge port configuration.**

This command displays the current bridge port configuration and flags.

**bridge fdb - forwarding database management**

**fdb** objects contain known Ethernet addresses on a link.

The corresponding commands display fdb entries, add new entries, append entries, and delete old ones.

**bridge fdb add - add a new fdb entry**

This command creates a new fdb entry.

**LLADDR**

the Ethernet MAC address.

**dev DEV**

the interface to which this address is associated.

**local** - is a local permanent fdb entry

**static** - is a static (no arp) fdb entry

**dynamic** - is a dynamic reachable age-able fdb entry

**self** - the address is associated with the port drivers fdb. Usually hardware.

**master** - the address is associated with master devices fdb. Usually software (default).

**router** - the destination address is associated with a router. Valid if the referenced device is a VXLAN type device and has route shortcircuit enabled.

**use** - the address is in use. User space can use this option to indicate to the kernel that the fdb entry is in use.

**extern\_learn** - this entry was learned externally. This option can be used to indicate to the kernel

that an entry was hardware or user-space controller learnt dynamic entry. Kernel will not age such an entry.

**sticky** - this entry will not change its port due to learning.

The next command line parameters apply only when the specified device *DEV* is of type VXLAN.

**dst IPADDR**

the IP address of the destination VXLAN tunnel endpoint where the Ethernet MAC ADDRESS resides.

**src\_vni SRC VNI**

the src VNI Network Identifier (or VXLAN Segment ID) this entry belongs to. Used only when the vxlan device is in external or collect metadata mode. If omitted the value specified at vxlan device creation will be used.

**vni VNI**

the VXLAN VNI Network Identifier (or VXLAN Segment ID) to use to connect to the remote VXLAN tunnel endpoint. If omitted the value specified at vxlan device creation will be used.

**port PORT**

the UDP destination PORT number to use to connect to the remote VXLAN tunnel endpoint. If omitted the default value is used.

**via DEVICE**

device name of the outgoing interface for the VXLAN device driver to reach the remote VXLAN tunnel endpoint.

**bridge fdb append - append a forwarding database entry**

This command adds a new fdb entry with an already known *LLADDR*. Valid only for multicast link layer addresses. The command adds support for broadcast and multicast Ethernet MAC addresses. The Ethernet MAC address is added multiple times into the forwarding database and the vxlan device driver sends a copy of the data packet to each entry found.

The arguments are the same as with **bridge fdb add**.

**bridge fdb delete - delete a forwarding database entry**

This command removes an existing fdb entry.

The arguments are the same as with **bridge fdb add**.

**bridge fdb replace - replace a forwarding database entry**

If no matching entry is found, a new one will be created instead.

The arguments are the same as with **bridge fdb add**.

**bridge fdb show - list forwarding entries.**

This command displays the current forwarding table.

With the **-statistics** option, the command becomes verbose. It prints out the last updated and last used time for each entry.

### **bridge mdb - multicast group database management**

**mdb** objects contain known IP multicast group addresses on a link.

The corresponding commands display mdb entries, add new entries, and delete old ones.

#### **bridge mdb add - add a new multicast group database entry**

This command creates a new mdb entry.

**dev** *DEV*

the interface where this group address is associated.

**port** *PORT*

the port whose link is known to have members of this multicast group.

**grp** *GROUP*

the IP multicast group address whose members reside on the link connected to the port.

**permanent** - the mdb entry is permanent

**temp** - the mdb entry is temporary (default)

**vid** *VID*

the VLAN ID which is known to have members of this multicast group.

#### **bridge mdb delete - delete a multicast group database entry**

This command removes an existing mdb entry.

The arguments are the same as with **bridge mdb add**.

#### **bridge mdb show - list multicast group database entries**

This command displays the current multicast group membership table. The table is populated by IGMP and MLD snooping in the bridge driver automatically. It can be altered by **bridge mdb add** and **bridge mdb del** commands manually too.

**dev** *DEV*

the interface only whose entries should be listed. Default is to list all bridge interfaces.

With the **-details** option, the command becomes verbose. It prints out the ports known to have a connected router.

With the **-statistics** option, the command displays timer values for mdb and router port entries.

**bridge vlan - VLAN filter list**

**vlan** objects contain known VLAN IDs for a link.

The corresponding commands display vlan filter entries, add new entries, and delete old ones.

**bridge vlan add - add a new vlan filter entry**

This command creates a new vlan filter entry.

**dev** *NAME*

the interface with which this vlan is associated.

**vid** *VID*

the VLAN ID that identifies the vlan.

**tunnel\_info** *TUNNEL\_ID*

the TUNNEL ID that maps to this vlan. The tunnel id is set in `dst_metadata` for every packet that belongs to this vlan (applicable to bridge ports with `vlan_tunnel` flag set).

**pvid** the vlan specified is to be considered a PVID at ingress. Any untagged frames will be assigned to this VLAN.

**untagged**

the vlan specified is to be treated as untagged on egress.

**self** the vlan is configured on the specified physical device. Required if the device is the bridge device.

**master** the vlan is configured on the software bridge (default).

**bridge vlan delete - delete a vlan filter entry**

This command removes an existing vlan filter entry.

The arguments are the same as with **bridge vlan add**. The **pvid** and **untagged** flags are ignored.

**bridge vlan show - list vlan configuration.**

This command displays the current VLAN filter table.

With the **-statistics** option, the command displays per-vlan traffic statistics.

**bridge vlan tunnelshow - list vlan tunnel mapping.**

This command displays the current vlan tunnel info mapping.

**bridge monitor - state monitoring**

The **bridge** utility can monitor the state of devices and addresses continuously. This option has a slightly different format. Namely, the **monitor** command is the first in the command line and then the object list follows:

**bridge monitor** [ **all** | *OBJECT-LIST* ]



*OBJECT-LIST* is the list of object types that we want to monitor. It may contain **link**, **fdb**, and **mdb**. If no **file** argument is given, **bridge** opens RTNETLINK, listens on it and dumps state changes in the format described in previous sections.

If a file name is given, it does not listen on RTNETLINK, but opens the file containing RTNETLINK messages saved in binary format and dumps them.

## NOTES

This command uses facilities added in Linux 3.0.

Although the forwarding table is maintained on a per-bridge device basis the bridge device is not part of the syntax. This is a limitation of the underlying netlink neighbour message protocol. When displaying the forwarding table, entries for all bridges are displayed. Add/delete/modify commands determine the underlying bridge device based on the bridge to which the corresponding ethernet device is attached.

## SEE ALSO

**ip(8)**

## BUGS

Please direct bugreports and patches to: [<netdev@vger.kernel.org>](mailto:netdev@vger.kernel.org)

## AUTHOR

Original Manpage by Stephen Hemminger