

NAME

`copy_file_range` – Copy a range of data from one file to another

SYNOPSIS

```
#define _GNU_SOURCE
#include <unistd.h>

ssize_t copy_file_range(int fd_in, loff_t *off_in,
                        int fd_out, loff_t *off_out,
                        size_t len, unsigned int flags);
```

DESCRIPTION

The `copy_file_range()` system call performs an in-kernel copy between two file descriptors without the additional cost of transferring data from the kernel to user space and then back into the kernel. It copies up to *len* bytes of data from file descriptor *fd_in* to file descriptor *fd_out*, overwriting any data that exists within the requested range of the target file.

The following semantics apply for *off_in*, and similar statements apply to *off_out*:

- * If *off_in* is NULL, then bytes are read from *fd_in* starting from the file offset, and the file offset is adjusted by the number of bytes copied.
- * If *off_in* is not NULL, then *off_in* must point to a buffer that specifies the starting offset where bytes from *fd_in* will be read. The file offset of *fd_in* is not changed, but *off_in* is adjusted appropriately.

The *flags* argument is provided to allow for future extensions and currently must be to 0.

RETURN VALUE

Upon successful completion, `copy_file_range()` will return the number of bytes copied between files. This could be less than the length originally requested.

On error, `copy_file_range()` returns `-1` and *errno* is set to indicate the error.

ERRORS**EBADF**

One or more file descriptors are not valid; or *fd_in* is not open for reading; or *fd_out* is not open for writing; or the **O_APPEND** flag is set for the open file description (see `open(2)`) referred to by the file descriptor *fd_out*.

EFBIG

An attempt was made to write a file that exceeds the implementation-defined maximum file size or the process's file size limit, or to write at a position past the maximum allowed offset.

EINVAL

Requested range extends beyond the end of the source file; or the *flags* argument is not 0.

EIO A low-level I/O error occurred while copying.

EISDIR

fd_in or *fd_out* refers to a directory.

ENOMEM

Out of memory.

ENOSPC

There is not enough space on the target filesystem to complete the copy.

EXDEV

The files referred to by *file_in* and *file_out* are not on the same mounted filesystem.

VERSIONS

The `copy_file_range()` system call first appeared in Linux 4.5, but glibc 2.27 provides a user-space emulation when it is not available.

CONFORMING TO

The **copy_file_range()** system call is a nonstandard Linux and GNU extension.

NOTES

If *file_in* is a sparse file, then **copy_file_range()** may expand any holes existing in the requested range. Users may benefit from calling **copy_file_range()** in a loop, and using the **lseek(2)** **SEEK_DATA** and **SEEK_HOLE** operations to find the locations of data segments.

copy_file_range() gives filesystems an opportunity to implement "copy acceleration" techniques, such as the use of reflinks (i.e., two or more inodes that share pointers to the same copy-on-write disk blocks) or server-side-copy (in the case of NFS).

EXAMPLE

```
#define _GNU_SOURCE
#include <fcntl.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <sys/syscall.h>
#include <unistd.h>

/* On versions of glibc before 2.27, we must invoke copy_file_range()
   using syscall(2) */

static loff_t
copy_file_range(int fd_in, loff_t *off_in, int fd_out,
                loff_t *off_out, size_t len, unsigned int flags)
{
    return syscall(__NR_copy_file_range, fd_in, off_in, fd_out,
                  off_out, len, flags);
}

int
main(int argc, char **argv)
{
    int fd_in, fd_out;
    struct stat stat;
    loff_t len, ret;

    if (argc != 3) {
        fprintf(stderr, "Usage: %s <source> <destination>\n", argv[0]);
        exit(EXIT_FAILURE);
    }

    fd_in = open(argv[1], O_RDONLY);
    if (fd_in == -1) {
        perror("open (argv[1])");
        exit(EXIT_FAILURE);
    }

    if (fstat(fd_in, &stat) == -1) {
        perror("fstat");
        exit(EXIT_FAILURE);
    }

    len = stat.st_size;
```

```
fd_out = open(argv[2], O_CREAT | O_WRONLY | O_TRUNC, 0644);
if (fd_out == -1) {
    perror("open (argv[2])");
    exit(EXIT_FAILURE);
}

do {
    ret = copy_file_range(fd_in, NULL, fd_out, NULL, len, 0);
    if (ret == -1) {
        perror("copy_file_range");
        exit(EXIT_FAILURE);
    }

    len -= ret;
} while (len > 0);

close(fd_in);
close(fd_out);
exit(EXIT_SUCCESS);
}
```

SEE ALSO

lseek(2), sendfile(2), splice(2)

COLOPHON

This page is part of release 5.02 of the Linux *man-pages* project. A description of the project, information about reporting bugs, and the latest version of this page, can be found at <https://www.kernel.org/doc/man-pages/>.