

**NAME**

ibv\_create\_qp\_ex, ibv\_destroy\_qp – create or destroy a queue pair (QP)

**SYNOPSIS**

```
#include <infiniband/verbs.h>
```

```
struct ibv_qp *ibv_create_qp_ex(struct ibv_context *context,
                                struct ibv_qp_init_attr_ex *qp_init_attr);
```

```
int ibv_destroy_qp(struct ibv_qp *qp);
```

**DESCRIPTION**

**ibv\_create\_qp\_ex()** creates a queue pair (QP) associated with the protection domain *pd*. The argument *qp\_init\_attr\_ex* is an *ibv\_qp\_init\_attr\_ex* struct, as defined in <infiniband/verbs.h>.

```
struct ibv_qp_init_attr_ex {
    void                *qp_context; /* Associated context of the QP */
    struct ibv_cq        *send_cq;   /* CQ to be associated with the Send Queue (SQ) */
    struct ibv_cq        *recv_cq;   /* CQ to be associated with the Receive Queue (RQ) */
    struct ibv_srq        *srq;      /* SRQ handle if QP is to be associated with an SRQ, otherwise NULL */
    struct ibv_qp_cap     cap;       /* QP capabilities */
    enum ibv_qp_type      qp_type;   /* QP Transport Service Type: IBV_QPT_RC, IBV_QPT_UC, IBV_QPT_GS */
    int                   sq_sig_all; /* If set, each Work Request (WR) submitted to the SQ generates a completion */
    uint32_t              comp_mask; /* Identifies valid fields */
    struct ibv_pd         *pd;       /* PD to be associated with the QP */
    struct ibv_xrca       *xrca;     /* XRC domain to be associated with the target QP */
    enum ibv_qp_create_flags create_flags; /* Creation flags for this QP */
    uint16_t              max_tso_header; /* Maximum TSO header size */
    struct ibv_rwrq_ind_table *rwrq_ind_tbl; /* Indirection table to be associated with the QP */
    struct ibv_rx_hash_conf rx_hash_conf; /* RX hash configuration to be used */
    uint32_t              source_qpn; /* Source QP number, creation flag IBV_QP_CREATE_SOURCE_QPN */
    uint64_t              send_ops_flags; /* Select which QP send ops will be defined in struct ibv_qp_ex. U */
};

struct ibv_qp_cap {
    uint32_t              max_send_wr; /* Requested max number of outstanding WRs in the SQ */
    uint32_t              max_recv_wr; /* Requested max number of outstanding WRs in the RQ */
    uint32_t              max_send_sge; /* Requested max number of scatter/gather (s/g) elements in a WR in the SQ */
    uint32_t              max_recv_sge; /* Requested max number of s/g elements in a WR in the RQ */
    uint32_t              max_inline_data; /* Requested max number of data (bytes) that can be posted inline */
};

enum ibv_qp_create_flags {
    IBV_QP_CREATE_BLOCK_SELF_MCAST_LB = 1 << 1, /* Prevent self multicast loopback */
    IBV_QP_CREATE_SCATTER_FCS = 1 << 8, /* FCS field will be scattered to host memory */
    IBV_QP_CREATE_CVLAN_STRIPPING = 1 << 9, /* CVLAN field will be stripped from incoming packets */
    IBV_QP_CREATE_SOURCE_QPN = 1 << 10, /* The created QP will use the source_qpn as its source */
    IBV_QP_CREATE_PCI_WRITE_END_PADDING = 1 << 11, /* Incoming packets will be padded to full packet size */
};

struct ibv_rx_hash_conf {
    uint8_t              rx_hash_function; /* RX hash function, use enum ibv_rx_hash_function_flags */
    uint8_t              rx_hash_key_len; /* RX hash key length */
    uint8_t              *rx_hash_key; /* RX hash key data */
    uint64_t             rx_hash_fields_mask; /* RX fields that should participate in the hashing, use enum ibv_rx_hash_fields */
};

enum ibv_rx_hash_fields {
    IBV_RX_HASH_SRC_IPV4 = 1 << 0,
```

```

        IBV_RX_HASH_DST_IPV4      = 1 << 1,
        IBV_RX_HASH_SRC_IPV6      = 1 << 2,
        IBV_RX_HASH_DST_IPV6      = 1 << 3,
        IBV_RX_HASH_SRC_PORT_TCP   = 1 << 4,
        IBV_RX_HASH_DST_PORT_TCP   = 1 << 5,
        IBV_RX_HASH_SRC_PORT_UDP   = 1 << 6,
        IBV_RX_HASH_DST_PORT_UDP   = 1 << 7,
        IBV_RX_HASH_IPSEC_SPI      = 1 << 8,
        /* When using tunneling protocol, e.g. VXLAN, then we have an inner (encapsulated packet) and outer
        * For applying RSS on the inner packet, then the following field should be set with one of the L3/L4
        */
        IBV_RX_HASH_INNER          = (1UL << 31),
};
struct ibv_qp_create_send_ops_flags {
        IBV_QP_EX_WITH_RDMA_WRITE      = 1 << 0,
        IBV_QP_EX_WITH_RDMA_WRITE_WITH_IMM = 1 << 1,
        IBV_QP_EX_WITH_SEND             = 1 << 2,
        IBV_QP_EX_WITH_SEND_WITH_IMM    = 1 << 3,
        IBV_QP_EX_WITH_RDMA_READ        = 1 << 4,
        IBV_QP_EX_WITH_ATOMIC_CMP_AND_SWP = 1 << 5,
        IBV_QP_EX_WITH_ATOMIC_FETCH_AND_ADD = 1 << 6,
        IBV_QP_EX_WITH_LOCAL_INV        = 1 << 7,
        IBV_QP_EX_WITH_BIND_MW          = 1 << 8,
        IBV_QP_EX_WITH_SEND_WITH_INV    = 1 << 9,
        IBV_QP_EX_WITH_TSO              = 1 << 10,
};

```

The function **ibv\_create\_qp\_ex()** will update the *qp\_init\_attr\_ex->cap* struct with the actual QP values of the QP that was created; the values will be greater than or equal to the values requested.

**ibv\_destroy\_qp()** destroys the QP *qp*.

## RETURN VALUE

**ibv\_create\_qp\_ex()** returns a pointer to the created QP, or NULL if the request fails. Check the QP number (**qp\_num**) in the returned QP.

**ibv\_destroy\_qp()** returns 0 on success, or the value of *errno* on failure (which indicates the failure reason).

## NOTES

The attributes *max\_rcv\_wr* and *max\_rcv\_sge* are ignored by **ibv\_create\_qp\_ex()** if the QP is to be associated with an SRQ.

The attribute *source\_qpn* is supported only on UD QP, without flow steering RX should not be possible.

Use **ibv\_qp\_to\_qp\_ex()** to get the *ibv\_qp\_ex* for accessing the send ops iterator interface, when QP create attr *IBV\_QP\_INIT\_ATTR\_SEND\_OPS\_FLAGS* is used.

**ibv\_destroy\_qp()** fails if the QP is attached to a multicast group.

**IBV\_QPT\_DRIVER** does not represent a specific service and is used for vendor specific QP logic.

## SEE ALSO

**ibv\_alloc\_pd(3)**, **ibv\_modify\_qp(3)**, **ibv\_query\_qp(3)**, **ibv\_create\_rmq\_ind\_table(3)**

## AUTHORS

Yishai Hadas <yishaih@mellanox.com>