

NAME

`ibv_post_send` – post a list of work requests (WRs) to a send queue

SYNOPSIS

```
#include <infiniband/verbs.h>
```

```
int ibv_post_send(struct ibv_qp *qp, struct ibv_send_wr *wr,
                  struct ibv_send_wr **bad_wr);
```

DESCRIPTION

ibv_post_send() posts the linked list of work requests (WRs) starting with *wr* to the send queue of the queue pair *qp*. It stops processing WRs from this list at the first failure (that can be detected immediately while requests are being posted), and returns this failing WR through *bad_wr*.

The argument *wr* is an `ibv_send_wr` struct, as defined in `<infiniband/verbs.h>`.

```
struct ibv_send_wr {
    uint64_t      wr_id;           /* User defined WR ID */
    struct ibv_send_wr *next;      /* Pointer to next WR in list, NULL if last WR */
    struct ibv_sge *sg_list;       /* Pointer to the s/g array */
    int           num_sge;         /* Size of the s/g array */
    enum ibv_wr_opcode opcode;     /* Operation type */
    int           send_flags;      /* Flags of the WR properties */
    union {
        __be32    imm_data;       /* Immediate data (in network byte order) */
        uint32_t   invalidate_rkey; /* Remote rkey to invalidate */
    };
    union {
        struct {
            uint64_t remote_addr; /* Start address of remote memory */
            uint32_t rkey;        /* Key of the remote Memory Region */
        } rdma;
        struct {
            uint64_t remote_addr; /* Start address of remote memory */
            uint64_t compare_add; /* Compare operand */
            uint64_t swap;        /* Swap operand */
            uint32_t rkey;        /* Key of the remote Memory Region */
        } atomic;
        struct {
            struct ibv_ah *ah;     /* Address handle (AH) for the remote QP */
            uint32_t remote_qpn;   /* QP number of the destination */
            uint32_t remote_qkey; /* Q_Key number of the destination */
        } ud;
    } wr;
    union {
        struct {
            uint32_t remote_srqn; /* Number of the remote SRQ */
        } xrc;
    } qp_type;
    union {
        struct {
            struct ibv_mw *mw;     /* Memory window (MW) */
            uint32_t rkey;        /* The desired new rkey of the MW */
            struct ibv_mw_bind_info bind_info; /* MW additional bind info */
        } bind_mw;
        struct {
            void *hdr;            /* Pointer address of inline data */
        } inline;
    };
};
```

```

                                uint16_t    hdr_sz; /* Inline header size */
                                uint16_t    mss;   /* Maximum segment size for each
                                } tso;
};

struct ibv_mw_bind_info {
    struct ibv_mr    *mr; /* The Memory region (MR) to bind the MW to */
    uint64_t        addr; /* The address the MW should start at */
    uint64_t        length; /* The length (in bytes) the MW should span */
    int             mw_access_flags; /* Access flags to the MW. Use ibv_access_flags */
};

struct ibv_sge {
    uint64_t        addr; /* Start address of the local memory buffer or number of bytes from
                                start of the MR for MRs which are IBV_ZERO_BASED */
    uint32_t        length; /* Length of the buffer */
    uint32_t        lkey; /* Key of the local Memory Region */
};

```

Each QP Transport Service Type supports a specific set of opcodes, as shown in the following table:

OPCODE	IBV_QPT_UD	IBV_QPT_UC	IBV_QPT_RC	IBV_QPT_XRC_SEND	IBV_QPT_RAW_PACKET
IBV_WR_SEND	X	X	X	X	X
IBV_WR_SEND_WITH_IMM	X	X	X	X	
IBV_WR_RDMA_WRITE		X	X	X	
IBV_WR_RDMA_WRITE_WITH_IMM			X	X	X
IBV_WR_RDMA_READ			X	X	
IBV_WR_ATOMIC_CMP_AND_SWP			X	X	
IBV_WR_ATOMIC_FETCH_AND_ADD			X	X	
IBV_WR_LOCAL_INV		X	X	X	
IBV_WR_BIND_MW		X	X	X	
IBV_WR_SEND_WITH_INV		X	X	X	
IBV_WR_TSO	X				X

The attribute `send_flags` describes the properties of the WR. It is either 0 or the bitwise OR of one or more of the following flags:

IBV_SEND_FENCE Set the fence indicator. Valid only for QPs with Transport Service Type **IBV_QPT_RC**

IBV_SEND_SIGNALED Set the completion notification indicator. Relevant only if QP was created with `sq_sig_all=0`

IBV_SEND_SOLICITED Set the solicited event indicator. Valid only for Send and RDMA Write with immediate

IBV_SEND_INLINE Send data in given gather list as inline data
in a send WQE. Valid only for Send and RDMA Write. The `L_Key` will not be checked.

IBV_SEND_IP_CSUM Offload the IPv4 and TCP/UDP checksum calculation.
Valid only when **device_cap_flags** in `device_attr` indicates current QP is supported by checksum offload.

RETURN VALUE

ibv_post_send() returns 0 on success, or the value of `errno` on failure (which indicates the failure reason).

NOTES

The user should not alter or destroy AHs associated with WRs until request is fully executed and a work completion has been retrieved from the corresponding completion queue (CQ) to avoid unexpected behavior.

The buffers used by a WR can only be safely reused after WR the request is fully executed and a work completion has been retrieved from the corresponding completion queue (CQ). However, if the IBV_SEND_INLINE flag was set, the buffer can be reused immediately after the call returns.

IBV_WR_DRIVER1 is an opcode that should be used to issue a specific driver operation.

SEE ALSO

ibv_create_qp(3), ibv_create_ah(3), ibv_post_recv(3), ibv_post_srq_recv(3), ibv_poll_cq(3)

AUTHORS

Dotan Barak <dotanba@gmail.com>

Majd Dibbiny <majd@mellanox.com>

Yishai Hadas <yishaih@mellanox.com>