

NAME

lwptut — An LWP Tutorial

DESCRIPTION

LWP (short for “Library for WWW in Perl”) is a very popular group of Perl modules for accessing data on the Web. Like most Perl module-distributions, each of LWP’s component modules comes with documentation that is a complete reference to its interface. However, there are so many modules in LWP that it’s hard to know where to start looking for information on how to do even the simplest most common things.

Really introducing you to using LWP would require a whole book — a book that just happens to exist, called *Perl & LWP*. But this article should give you a taste of how you can go about some common tasks with LWP.

Getting documents with LWP::Simple

If you just want to get what’s at a particular URL, the simplest way to do it is LWP::Simple’s functions.

In a Perl program, you can call its `get($url)` function. It will try getting that URL’s content. If it works, then it’ll return the content; but if there’s some error, it’ll return `undef`.

```
my $url = 'http://www.npr.org/programs/fa/?todayDate=current';
# Just an example: the URL for the most recent /Fresh Air/ show

use LWP::Simple;
my $content = get $url;
die "Couldn't get $url" unless defined $content;

# Then go do things with $content, like this:

if($content =~ m/jazz/i) {
    print "They're talking about jazz today on Fresh Air!\n";
}
else {
    print "Fresh Air is apparently jazzless today.\n";
}
```

The handiest variant on `get` is `getprint`, which is useful in Perl one-liners. If it can get the page whose URL you provide, it sends it to `STDOUT`; otherwise it complains to `STDERR`.

```
% perl -MLWP::Simple -e "getprint 'http://www.cpan.org/RECENT'"
```

That is the URL of a plain text file that lists new files in CPAN in the past two weeks. You can easily make it part of a tidy little shell command, like this one that mails you the list of new Acme:: modules:

```
% perl -MLWP::Simple -e "getprint 'http://www.cpan.org/RECENT' " \
| grep "/by-module/Acme" | mail -s "New Acme modules! Joy!" $USER
```

There are other useful functions in LWP::Simple, including one function for running a HEAD request on a URL (useful for checking links, or getting the last-revised time of a URL), and two functions for saving/mirroring a URL to a local file. See the LWP::Simple documentation for the full details, or chapter 2 of *Perl & LWP* for more examples.

The Basics of the LWP Class Model

LWP::Simple’s functions are handy for simple cases, but its functions don’t support cookies or authorization, don’t support setting header lines in the HTTP request, generally don’t support reading header lines in the HTTP response (notably the full HTTP error message, in case of an error). To get at all those features, you’ll have to use the full LWP class model.

While LWP consists of dozens of classes, the main two that you have to understand are LWP::UserAgent and HTTP::Response. LWP::UserAgent is a class for “virtual browsers” which you use for performing requests, and HTTP::Response is a class for the responses (or error messages) that you get back from those requests.

The basic idiom is `$response = $browser->get($url)`, or more fully illustrated:

```
# Early in your program:

use LWP 5.64; # Loads all important LWP classes, and makes
              # sure your version is reasonably recent.

my $browser = LWP::UserAgent->new;

...

# Then later, whenever you need to make a get request:
my $url = 'http://www.npr.org/programs/fa/?todayDate=current';

my $response = $browser->get( $url );
die "Can't get $url -- ", $response->status_line
    unless $response->is_success;

die "Hey, I was expecting HTML, not ", $response->content_type
    unless $response->content_type eq 'text/html';
    # or whatever content-type you're equipped to deal with

# Otherwise, process the content somehow:

if($response->decoded_content =~ m/jazz/i) {
    print "They're talking about jazz today on Fresh Air!\n";
}
else {
    print "Fresh Air is apparently jazzless today.\n";
}
```

There are two objects involved: `$browser`, which holds an object of class `LWP::UserAgent`, and then the `$response` object, which is of class `HTTP::Response`. You really need only one browser object per program; but every time you make a request, you get back a new `HTTP::Response` object, which will have some interesting attributes:

- A status code indicating success or failure (which you can test with `$response->is_success`).
- An HTTP status line that is hopefully informative if there's failure (which you can see with `$response->status_line`, returning something like "404 Not Found").
- A MIME content-type like "text/html", "image/gif", "application/xml", etc., which you can see with `$response->content_type`
- The actual content of the response, in `$response->decoded_content`. If the response is HTML, that's where the HTML source will be; if it's a GIF, then `$response->decoded_content` will be the binary GIF data.
- And dozens of other convenient and more specific methods that are documented in the docs for `HTTP::Response`, and its superclasses `HTTP::Message` and `HTTP::Headers`.

Adding Other HTTP Request Headers

The most commonly used syntax for requests is `$response = $browser->get($url)`, but in truth, you can add extra HTTP header lines to the request by adding a list of key-value pairs after the URL, like so:

```
$response = $browser->get( $url, $key1, $value1, $key2, $value2, ... );
```

For example, here's how to send some commonly used headers, in case you're dealing with a site that would otherwise reject your request:

```
my @ns_headers = (
    'User-Agent' => 'Mozilla/4.76 [en] (Win98; U)',
    'Accept' => 'image/gif, image/x-xbitmap, image/jpeg, image/pjpeg, image/png, *',
    'Accept-Charset' => 'iso-8859-1,*,utf-8',
    'Accept-Language' => 'en-US',
);

...
```

```
$response = $browser->get($url, @ns_headers);
```

If you weren't reusing that array, you could just go ahead and do this:

```
$response = $browser->get($url,
    'User-Agent' => 'Mozilla/4.76 [en] (Win98; U)',
    'Accept' => 'image/gif, image/x-xbitmap, image/jpeg, image/pjpeg, image/png, *',
    'Accept-Charset' => 'iso-8859-1,*,utf-8',
    'Accept-Language' => 'en-US',
);
```

If you were only ever changing the 'User-Agent' line, you could just change the \$browser object's default line from "libwww-perl/5.65" (or the like) to whatever you like, using the LWP::UserAgent agent method:

```
$browser->agent('Mozilla/4.76 [en] (Win98; U)');
```

Enabling Cookies

A default LWP::UserAgent object acts like a browser with its cookies support turned off. There are various ways of turning it on, by setting its `cookie_jar` attribute. A "cookie jar" is an object representing a little database of all the HTTP cookies that a browser knows about. It can correspond to a file on disk or an in-memory object that starts out empty, and whose collection of cookies will disappear once the program is finished running.

To give a browser an in-memory empty cookie jar, you set its `cookie_jar` attribute like so:

```
use HTTP::CookieJar::LWP;
$browser->cookie_jar( HTTP::CookieJar::LWP->new );
```

To save a cookie jar to disk, see "dump_cookies" in HTTP::CookieJar. To load cookies from disk into a jar, see "load_cookies" in HTTP::CookieJar.

Posting Form Data

Many HTML forms send data to their server using an HTTP POST request, which you can send with this syntax:

```
$response = $browser->post( $url,
    [
        formkey1 => value1,
        formkey2 => value2,
        ...
    ],
);
```

Or if you need to send HTTP headers:

```

$response = $browser->post( $url,
    [
        formkey1 => value1,
        formkey2 => value2,
        ...
    ],
    headerkey1 => value1,
    headerkey2 => value2,
);

```

For example, the following program makes a search request to AltaVista (by sending some form data via an HTTP POST request), and extracts from the HTML the report of the number of matches:

```

use strict;
use warnings;
use LWP 5.64;
my $browser = LWP::UserAgent->new;

my $word = 'tarragon';

my $url = 'http://search.yahoo.com/yhs/search';
my $response = $browser->post( $url,
    [ 'q' => $word, # the Altavista query string
      'fr' => 'altavista', 'pg' => 'q', 'avkw' => 'tgz', 'kl' => 'XX',
    ]
);
die "$url error: ", $response->status_line
    unless $response->is_success;
die "Weird content type at $url -- ", $response->content_type
    unless $response->content_is_html;

if( $response->decoded_content =~ m{([0-9,]+)(?:<.*?>)? results for} ) {
    # The substring will be like "996,000</strong> results for"
    print "$word: $1\n";
}
else {
    print "Couldn't find the match-string in the response\n";
}

```

Sending GET Form Data

Some HTML forms convey their form data not by sending the data in an HTTP POST request, but by making a normal GET request with the data stuck on the end of the URL. For example, if you went to www.imdb.com and ran a search on “Blade Runner”, the URL you’d see in your browser window would be:

```
http://www.imdb.com/find?s=all&q=Blade+Runner
```

To run the same search with LWP, you’d use this idiom, which involves the URI class:

```

use URI;
my $url = URI->new( 'http://www.imdb.com/find' );
    # makes an object representing the URL

$url->query_form( # And here the form data pairs:
    'q' => 'Blade Runner',
    's' => 'all',
);

```

```
my $response = $browser->get($url);
```

See chapter 5 of *Perl & LWP* for a longer discussion of HTML forms and of form data, and chapters 6 through 9 for a longer discussion of extracting data from HTML.

Absolutizing URLs

The URI class that we just mentioned above provides all sorts of methods for accessing and modifying parts of URLs (such as asking sort of URL it is with `$url->scheme`, and asking what host it refers to with `$url->host`, and so on, as described in the docs for the URI class. However, the methods of most immediate interest are the `query_form` method seen above, and now the `new_abs` method for taking a probably-relative URL string (like `“./foo.html”`) and getting back an absolute URL (like `“http://www.perl.com/stuff/foo.html”`), as shown here:

```
use URI;
$abs = URI->new_abs($maybe_relative, $base);
```

For example, consider this program that matches URLs in the HTML list of new modules in CPAN:

```
use strict;
use warnings;
use LWP;
my $browser = LWP::UserAgent->new;

my $url = 'http://www.cpan.org/RECENT.html';
my $response = $browser->get($url);
die "Can't get $url -- ", $response->status_line
    unless $response->is_success;

my $html = $response->decoded_content;
while( $html =~ m/<A HREF=\"(.*)\"/g ) {
    print "$1\n";
}
```

When run, it emits output that starts out something like this:

```
MIRRORING.FROM
RECENT
RECENT.html
authors/00whois.html
authors/01mailrc.txt.gz
authors/id/A/AA/AASSAD/CHECKSUMS
...
```

However, if you actually want to have those be absolute URLs, you can use the URI module's `new_abs` method, by changing the while loop to this:

```
while( $html =~ m/<A HREF=\"(.*)\"/g ) {
    print URI->new_abs( $1, $response->base ) , "\n";
}
```

(The `$response->base` method from `HTTP::Message` is for returning what URL should be used for resolving relative URLs — it's usually just the same as the URL that you requested.)

That program then emits nicely absolute URLs:

```

http://www.cpan.org/MIRRORING.FROM
http://www.cpan.org/RECENT
http://www.cpan.org/RECENT.html
http://www.cpan.org/authors/00whois.html
http://www.cpan.org/authors/01mailrc.txt.gz
http://www.cpan.org/authors/id/A/AA/AASSAD/CHECKSUMS
...

```

See chapter 4 of *Perl & LWP* for a longer discussion of URI objects.

Of course, using a regexp to match hrefs is a bit simplistic, and for more robust programs, you'll probably want to use an HTML-parsing module like `HTML::LinkExtor` or `HTML::Tokenizer` or even maybe `HTML::TreeBuilder`.

Other Browser Attributes

`LWP::UserAgent` objects have many attributes for controlling how they work. Here are a few notable ones:

- `$browser->timeout(15);`

This sets this browser object to give up on requests that don't answer within 15 seconds.

- `$browser->protocols_allowed(['http', 'gopher']);`

This sets this browser object to not speak any protocols other than HTTP and gopher. If it tries accessing any other kind of URL (like an "ftp:" or "mailto:" or "news:" URL), then it won't actually try connecting, but instead will immediately return an error code 500, with a message like "Access to 'ftp' URIs has been disabled".

- `use LWP::ConnCache; $browser->conn_cache(LWP::ConnCache->new());`

This tells the browser object to try using the HTTP/1.1 "Keep-Alive" feature, which speeds up requests by reusing the same socket connection for multiple requests to the same server.

- `$browser->agent('SomeName/1.23 (more info here maybe)');`

This changes how the browser object will identify itself in the default "User-Agent" line in its HTTP requests. By default, it'll send "libwww-perl/versionnumber", like "libwww-perl/5.65". You can change that to something more descriptive like this:

```
$browser->agent( 'SomeName/3.14 (contact@robotplexus.int)' );
```

Or if need be, you can go in disguise, like this:

```
$browser->agent( 'Mozilla/4.0 (compatible; MSIE 5.12; Mac_PowerPC)' );
```

- `push @{$browser->requests_redirectable}, 'POST';`

This tells this browser to obey redirection responses to POST requests (like most modern interactive browsers), even though the HTTP RFC says that should not normally be done.

For more options and information, see the full documentation for `LWP::UserAgent`.

Writing Polite Robots

If you want to make sure that your LWP-based program respects *robots.txt* files and doesn't make too many requests too fast, you can use the `LWP::RobotUA` class instead of the `LWP::UserAgent` class.

`LWP::RobotUA` class is just like `LWP::UserAgent`, and you can use it like so:

```

use LWP::RobotUA;
my $browser = LWP::RobotUA->new('YourSuperBot/1.34', 'you@yoursite.com');
# Your bot's name and your email address

my $response = $browser->get($url);

```

But `LWP::RobotUA` adds these features:

- If the *robots.txt* on *\$url*'s server forbids you from accessing *\$url*, then the *\$browser* object (assuming it's of class *LWP::RobotUA*) won't actually request it, but instead will give you back (in *\$response*) a 403 error with a message "Forbidden by robots.txt". That is, if you have this line:

```
die "$url -- ", $response->status_line, "\nAborted"
unless $response->is_success;
```

then the program would die with an error message like this:

```
http://whatever.site.int/pith/x.html -- 403 Forbidden by robots.txt
Aborted at whateverprogram.pl line 1234
```

- If this *\$browser* object sees that the last time it talked to *\$url*'s server was too recently, then it will pause (via *sleep*) to avoid making too many requests too often. How long it will pause for, is by default one minute — but you can control it with the *\$browser->delay(minutes)* attribute.

For example, this code:

```
$browser->delay( 7/60 );
```

...means that this browser will pause when it needs to avoid talking to any given server more than once every 7 seconds.

For more options and information, see the full documentation for *LWP::RobotUA*.

Using Proxies

In some cases, you will want to (or will have to) use proxies for accessing certain sites and/or using certain protocols. This is most commonly the case when your LWP program is running (or could be running) on a machine that is behind a firewall.

To make a browser object use proxies that are defined in the usual environment variables (*HTTP_PROXY*, etc.), just call the *env_proxy* on a user-agent object before you go making any requests on it. Specifically:

```
use LWP::UserAgent;
my $browser = LWP::UserAgent->new;

# And before you go making any requests:
$browser->env_proxy;
```

For more information on proxy parameters, see the *LWP::UserAgent* documentation, specifically the *proxy*, *env_proxy*, and *no_proxy* methods.

HTTP Authentication

Many web sites restrict access to documents by using "HTTP Authentication". This isn't just any form of "enter your password" restriction, but is a specific mechanism where the HTTP server sends the browser an HTTP code that says "That document is part of a protected 'realm', and you can access it only if you re-request it and add some special authorization headers to your request".

For example, the Unicode.org admins stop email-harvesting bots from harvesting the contents of their mailing list archives, by protecting them with HTTP Authentication, and then publicly stating the username and password (at <http://www.unicode.org/mail-arch/>) — namely username "unicode-ml" and password "unicode".

For example, consider this URL, which is part of the protected area of the web site:

```
http://www.unicode.org/mail-arch/unicode-ml/y2002-m08/0067.html
```

If you access that with a browser, you'll get a prompt like "Enter username and password for 'Unicode-MailList-Archives' at server 'www.unicode.org'".

In LWP, if you just request that URL, like this:

```

use LWP;
my $browser = LWP::UserAgent->new;

my $url =
    'http://www.unicode.org/mail-arch/unicode-ml/y2002-m08/0067.html';
my $response = $browser->get($url);

die "Error: ", $response->header('WWW-Authenticate') || 'Error accessing',
    # ('WWW-Authenticate' is the realm-name)
    "\n ", $response->status_line, "\n at $url\n Aborting"
    unless $response->is_success;

```

Then you'll get this error:

```

Error: Basic realm="Unicode-MailList-Archives"
401 Authorization Required
at http://www.unicode.org/mail-arch/unicode-ml/y2002-m08/0067.html
Aborting at auth1.pl line 9. [or wherever]

```

...because the `$browser` doesn't know any the username and password for that realm ("Unicode-MailList-Archives") at that host ("www.unicode.org"). The simplest way to let the browser know about this is to use the `credentials` method to let it know about a username and password that it can try using for that realm at that host. The syntax is:

```

$browser->credentials(
    'servername:portnumber',
    'realm-name',
    'username' => 'password'
);

```

In most cases, the port number is 80, the default TCP/IP port for HTTP; and you usually call the `credentials` method before you make any requests. For example:

```

$browser->credentials(
    'reports.mybazouki.com:80',
    'web_server_usage_reports',
    'plinky' => 'banjo123'
);

```

So if we add the following to the program above, right after the `$browser = LWP::UserAgent->new;` line...

```

$browser->credentials( # add this to our $browser 's "key ring"
    'www.unicode.org:80',
    'Unicode-MailList-Archives',
    'unicode-ml' => 'unicode'
);

```

...then when we run it, the request succeeds, instead of causing the `die` to be called.

Accessing HTTPS URLs

When you access an HTTPS URL, it'll work for you just like an HTTP URL would — if your LWP installation has HTTPS support (via an appropriate Secure Sockets Layer library). For example:


```

use LWP;
my $url = 'https://www.paypal.com/';    # Yes, HTTPS!
my $browser = LWP::UserAgent->new;
my $response = $browser->get($url);
die "Error at $url\n ", $response->status_line, "\n Aborting"
    unless $response->is_success;
print "Whee, it worked!  I got that ",
    $response->content_type, " document!\n";

```

If your LWP installation doesn't have HTTPS support set up, then the response will be unsuccessful, and you'll get this error message:

```

Error at https://www.paypal.com/
501 Protocol scheme 'https' is not supported
Aborting at paypal.pl line 7.    [or whatever program and line]

```

If your LWP installation *does* have HTTPS support installed, then the response should be successful, and you should be able to consult `$response` just like with any normal HTTP response.

For information about installing HTTPS support for your LWP installation, see the helpful *README.SSL* file that comes in the libwww-perl distribution.

Getting Large Documents

When you're requesting a large (or at least potentially large) document, a problem with the normal way of using the request methods (like `$response = $browser->get($url)`) is that the response object in memory will have to hold the whole document — *in memory*. If the response is a thirty megabyte file, this is likely to be quite an imposition on this process's memory usage.

A notable alternative is to have LWP save the content to a file on disk, instead of saving it up in memory. This is the syntax to use:

```

$response = $ua->get($url,
    ':content_file' => $filespec,
);

```

For example,

```

$response = $ua->get('http://search.cpan.org/',
    ':content_file' => '/tmp/sco.html'
);

```

When you use this `:content_file` option, the `$response` will have all the normal header lines, but `$response->content` will be empty. Errors writing to the content file (for example due to permission denied or the filesystem being full) will be reported via the `Client-Aborted` or `X-Died` response headers, and not the `is_success` method:

```

if ($response->header('Client-Aborted') eq 'die') {
    # handle error ...
}

```

Note that this `:content_file` option isn't supported under older versions of LWP, so you should consider adding `use LWP 5.66;` to check the LWP version, if you think your program might run on systems with older versions.

If you need to be compatible with older LWP versions, then use this syntax, which does the same thing:

```

use HTTP::Request::Common;
$response = $ua->request( GET($url), $filespec );

```

SEE ALSO

Remember, this article is just the most rudimentary introduction to LWP — to learn more about LWP and LWP-related tasks, you really must read from the following:

- `LWP::Simple` — simple functions for getting/heading/mirroring URLs

- LWP — overview of the libwww-perl modules
- LWP::UserAgent — the class for objects that represent “virtual browsers”
- HTTP::Response — the class for objects that represent the response to a LWP response, as in `$response = $browser->get(...)`
- HTTP::Message and HTTP::Headers — classes that provide more methods to HTTP::Response.
- URI — class for objects that represent absolute or relative URLs
- URI::Escape — functions for URL-escaping and URL-unescaping strings (like turning “this & that” to and from “this%20%26%20that”).
- HTML::Entities — functions for HTML-escaping and HTML-unescaping strings (like turning “C. & E. Brontë” to and from “C. & E. Brontë”).
- HTML::TokeParser and HTML::TreeBuilder — classes for parsing HTML
- HTML::LinkExtor — class for finding links in HTML documents
- The book *Perl & LWP* by Sean M. Burke. O’Reilly & Associates, 2002. ISBN: 0-596-00178-9, <<http://oreilly.com/catalog/perlwp/>>. The whole book is also available free online: <<http://lwp.interglacial.com/>>.

COPYRIGHT

Copyright 2002, Sean M. Burke. You can redistribute this document and/or modify it, but only under the same terms as Perl itself.

AUTHOR

Sean M. Burke sburke@cpan.org