

NAME

lockf – apply, test or remove a POSIX lock on an open file

SYNOPSIS

```
#include <unistd.h>
```

```
int lockf(int fd, int cmd, off_t len);
```

Feature Test Macro Requirements for glibc (see **feature_test_macros(7)**):

```
lockf():
```

```
    _XOPEN_SOURCE >= 500
```

```
    || /* Glibc since 2.19: */ _DEFAULT_SOURCE
```

```
    || /* Glibc versions <= 2.19: */ _BSD_SOURCE || _SVID_SOURCE
```

DESCRIPTION

Apply, test or remove a POSIX lock on a section of an open file. The file is specified by *fd*, a file descriptor open for writing, the action by *cmd*, and the section consists of byte positions *pos*..*pos+len-1* if *len* is positive, and *pos-len*..*pos-1* if *len* is negative, where *pos* is the current file position, and if *len* is zero, the section extends from the current file position to infinity, encompassing the present and future end-of-file positions. In all cases, the section may extend past current end-of-file.

On Linux, **lockf()** is just an interface on top of **fcntl(2)** locking. Many other systems implement **lockf()** in this way, but note that POSIX.1 leaves the relationship between **lockf()** and **fcntl(2)** locks unspecified. A portable application should probably avoid mixing calls to these interfaces.

Valid operations are given below:

F_LOCK

Set an exclusive lock on the specified section of the file. If (part of) this section is already locked, the call blocks until the previous lock is released. If this section overlaps an earlier locked section, both are merged. File locks are released as soon as the process holding the locks closes some file descriptor for the file. A child process does not inherit these locks.

F_TLOCK

Same as **F_LOCK** but the call never blocks and returns an error instead if the file is already locked.

F_ULOCK

Unlock the indicated section of the file. This may cause a locked section to be split into two locked sections.

F_TEST

Test the lock: return 0 if the specified section is unlocked or locked by this process; return -1, set *errno* to **EAGAIN** (**EACCES** on some other systems), if another process holds a lock.

RETURN VALUE

On success, zero is returned. On error, -1 is returned, and *errno* is set appropriately.

ERRORS**EACCES** or **EAGAIN**

The file is locked and **F_TLOCK** or **F_TEST** was specified, or the operation is prohibited because the file has been memory-mapped by another process.

EBADF

fd is not an open file descriptor; or *cmd* is **F_LOCK** or **F_TLOCK** and *fd* is not a writable file descriptor.

EDEADLK

The command was **F_LOCK** and this lock operation would cause a deadlock.

EINTR

While waiting to acquire a lock, the call was interrupted by delivery of a signal caught by a handler; see **signal(7)**.

EINVAL

An invalid operation was specified in *cmd*.

ENOLCK

Too many segment locks open, lock table is full.

ATTRIBUTES

For an explanation of the terms used in this section, see **attributes(7)**.

Interface	Attribute	Value
lockf()	Thread safety	MT-Safe

CONFORMING TO

POSIX.1-2001, POSIX.1-2008, SVr4.

SEE ALSO

fcntl(2), **flock(2)**

locks.txt and *mandatory-locking.txt* in the Linux kernel source directory *Documentation/filesystems* (on older kernels, these files are directly under the *Documentation* directory, and *mandatory-locking.txt* is called *mandatory.txt*)

COLOPHON

This page is part of release 5.02 of the Linux *man-pages* project. A description of the project, information about reporting bugs, and the latest version of this page, can be found at <https://www.kernel.org/doc/man-pages/>.