

NAME

sendfile – transfer data between file descriptors

SYNOPSIS

```
#include <sys/sendfile.h>
```

```
ssize_t sendfile(int out_fd, int in_fd, off_t *offset, size_t count);
```

DESCRIPTION

sendfile() copies data between one file descriptor and another. Because this copying is done within the kernel, **sendfile()** is more efficient than the combination of **read(2)** and **write(2)**, which would require transferring data to and from user space.

in_fd should be a file descriptor opened for reading and *out_fd* should be a descriptor opened for writing.

If *offset* is not NULL, then it points to a variable holding the file offset from which **sendfile()** will start reading data from *in_fd*. When **sendfile()** returns, this variable will be set to the offset of the byte following the last byte that was read. If *offset* is not NULL, then **sendfile()** does not modify the file offset of *in_fd*; otherwise the file offset is adjusted to reflect the number of bytes read from *in_fd*.

If *offset* is NULL, then data will be read from *in_fd* starting at the file offset, and the file offset will be updated by the call.

count is the number of bytes to copy between the file descriptors.

The *in_fd* argument must correspond to a file which supports **mmap(2)**-like operations (i.e., it cannot be a socket).

In Linux kernels before 2.6.33, *out_fd* must refer to a socket. Since Linux 2.6.33 it can be any file. If it is a regular file, then **sendfile()** changes the file offset appropriately.

RETURN VALUE

If the transfer was successful, the number of bytes written to *out_fd* is returned. Note that a successful call to **sendfile()** may write fewer bytes than requested; the caller should be prepared to retry the call if there were unsent bytes. See also NOTES.

On error, -1 is returned, and *errno* is set appropriately.

ERRORS**EAGAIN**

Nonblocking I/O has been selected using **O_NONBLOCK** and the write would block.

EBADF

The input file was not opened for reading or the output file was not opened for writing.

EFAULT

Bad address.

EINVAL

Descriptor is not valid or locked, or an **mmap(2)**-like operation is not available for *in_fd*, or *count* is negative.

EINVAL

out_fd has the **O_APPEND** flag set. This is not currently supported by **sendfile()**.

EIO

Unspecified error while reading from *in_fd*.

ENOMEM

Insufficient memory to read from *in_fd*.

EOVERFLOW

count is too large, the operation would result in exceeding the maximum size of either the input file or the output file.

ESPIPE

offset is not NULL but the input file is not **seek(2)**-able.

VERSIONS

sendfile() first appeared in Linux 2.2. The include file `<sys/sendfile.h>` is present since glibc 2.1.

CONFORMING TO

Not specified in POSIX.1-2001, nor in other standards.

Other UNIX systems implement **sendfile()** with different semantics and prototypes. It should not be used in portable programs.

NOTES

sendfile() will transfer at most 0x7ffff000 (2,147,479,552) bytes, returning the number of bytes actually transferred. (This is true on both 32-bit and 64-bit systems.)

If you plan to use **sendfile()** for sending files to a TCP socket, but need to send some header data in front of the file contents, you will find it useful to employ the **TCP_CORK** option, described in **tcp(7)**, to minimize the number of packets and to tune performance.

In Linux 2.4 and earlier, *out_fd* could also refer to a regular file; this possibility went away in the Linux 2.6.x kernel series, but was restored in Linux 2.6.33.

The original Linux **sendfile()** system call was not designed to handle large file offsets. Consequently, Linux 2.4 added **sendfile64()**, with a wider type for the *offset* argument. The glibc **sendfile()** wrapper function transparently deals with the kernel differences.

Applications may wish to fall back to **read(2)/write(2)** in the case where **sendfile()** fails with **EINVAL** or **ENOSYS**.

If *out_fd* refers to a socket or pipe with zero-copy support, callers must ensure the transferred portions of the file referred to by *in_fd* remain unmodified until the reader on the other end of *out_fd* has consumed the transferred data.

The Linux-specific **splice(2)** call supports transferring data between arbitrary file descriptors provided one (or both) of them is a pipe.

SEE ALSO

copy_file_range(2), **mmap(2)**, **open(2)**, **socket(2)**, **splice(2)**

COLOPHON

This page is part of release 5.02 of the Linux *man-pages* project. A description of the project, information about reporting bugs, and the latest version of this page, can be found at <https://www.kernel.org/doc/man-pages/>.