

**NAME**

**sprof** – read and display shared object profiling data

**SYNOPSIS**

**sprof** [*option*]... *shared-object-path* [*profile-data-path*]

**DESCRIPTION**

The **sprof** command displays a profiling summary for the shared object (shared library) specified as its first command-line argument. The profiling summary is created using previously generated profiling data in the (optional) second command-line argument. If the profiling data pathname is omitted, then **sprof** will attempt to deduce it using the soname of the shared object, looking for a file with the name *<soname>.profile* in the current directory.

**OPTIONS**

The following command-line options specify the profile output to be produced:

**-c, --call-pairs**

Print a list of pairs of call paths for the interfaces exported by the shared object, along with the number of times each path is used.

**-p, --flat-profile**

Generate a flat profile of all of the functions in the monitored object, with counts and ticks.

**-q, --graph**

Generate a call graph.

If none of the above options is specified, then the default behavior is to display a flat profile and a call graph.

The following additional command-line options are available:

**–?, --help**

Display a summary of command-line options and arguments and exit.

**--usage**

Display a short usage message and exit.

**-V, --version**

Display the program version and exit.

**CONFORMING TO**

The **sprof** command is a GNU extension, not present in POSIX.1.

**EXAMPLE**

The following example demonstrates the use of **sprof**. The example consists of a main program that calls two functions in a shared object. First, the code of the main program:

```
$ cat prog.c
#include <stdlib.h>

void x1(void);
void x2(void);

int
main(int argc, char *argv[])
{
    x1 ();
    x2 ();
    exit (EXIT_SUCCESS);
}
```

The functions *x1()* and *x2()* are defined in the following source file that is used to construct the shared object:

```

$ cat libdemo.c
#include <unistd.h>

void
consumeCpu1(int lim)
{
    int j;

    for (j = 0; j < lim; j++)
        getppid();
}

void
x1(void) {
    int j;

    for (j = 0; j < 100; j++)
        consumeCpu1(200000);
}

void
consumeCpu2(int lim)
{
    int j;

    for (j = 0; j < lim; j++)
        getppid();
}

void
x2(void)
{
    int j;

    for (j = 0; j < 1000; j++)
        consumeCpu2(10000);
}

```

Now we construct the shared object with the real name *libdemo.so.1.0.1*, and the soname *libdemo.so.1*:

```

$ cc -g -fPIC -shared -Wl,-soname,libdemo.so.1 \
  -o libdemo.so.1.0.1 libdemo.c

```

Then we construct symbolic links for the library soname and the library linker name:

```

$ ln -sf libdemo.so.1.0.1 libdemo.so.1
$ ln -sf libdemo.so.1 libdemo.so

```

Next, we compile the main program, linking it against the shared object, and then list the dynamic dependencies of the program:

```

$ cc -g -o prog prog.c -L. -ldemo
$ ldd prog
    linux-vdso.so.1 => (0x00007fff86d66000)
    libdemo.so.1 => not found
    libc.so.6 => /lib64/libc.so.6 (0x00007fd4dc138000)
    /lib64/ld-linux-x86-64.so.2 (0x00007fd4dc51f000)

```

In order to get profiling information for the shared object, we define the environment variable

**LD\_PROFILE** with the soname of the library:

```
$ export LD_PROFILE=libdemo.so.1
```

We then define the environment variable **LD\_PROFILE\_OUTPUT** with the pathname of the directory where profile output should be written, and create that directory if it does not exist already:

```
$ export LD_PROFILE_OUTPUT=$(pwd)/prof_data
$ mkdir -p $LD_PROFILE_OUTPUT
```

**LD\_PROFILE** causes profiling output to be *appended* to the output file if it already exists, so we ensure that there is no preexisting profiling data:

```
$ rm -f $LD_PROFILE_OUTPUT/$LD_PROFILE.profile
```

We then run the program to produce the profiling output, which is written to a file in the directory specified in **LD\_PROFILE\_OUTPUT**:

```
$ LD_LIBRARY_PATH=. ./prog
$ ls prof_data
libdemo.so.1.profile
```

We then use the **sprof -p** option to generate a flat profile with counts and ticks:

```
$ sprof -p libdemo.so.1 $LD_PROFILE_OUTPUT/libdemo.so.1.profile
Flat profile:
```

Each sample counts as 0.01 seconds.

%	cumulative	self		self	total	
time	seconds	seconds	calls	us/call	us/call	name
60.00	0.06	0.06	100	600.00		consumeCpu1
40.00	0.10	0.04	1000	40.00		consumeCpu2
0.00	0.10	0.00	1	0.00		x1
0.00	0.10	0.00	1	0.00		x2

The **sprof -q** option generates a call graph:

```
$ sprof -q libdemo.so.1 $LD_PROFILE_OUTPUT/libdemo.so.1.profile
```

index	% time	self	children	called	name
		0.00	0.00	100/100	x1 [1]
[0]	100.0	0.00	0.00	100	consumeCpu1 [0]
-----					
		0.00	0.00	1/1	<UNKNOWN>
[1]	0.0	0.00	0.00	1	x1 [1]
		0.00	0.00	100/100	consumeCpu1 [0]
-----					
		0.00	0.00	1000/1000	x2 [3]
[2]	0.0	0.00	0.00	1000	consumeCpu2 [2]
-----					
		0.00	0.00	1/1	<UNKNOWN>
[3]	0.0	0.00	0.00	1	x2 [3]
		0.00	0.00	1000/1000	consumeCpu2 [2]
-----					

Above and below, the "<UNKNOWN>" strings represent identifiers that are outside of the profiled object (in this example, these are instances of *main()*).

The **sprof -c** option generates a list of call pairs and the number of their occurrences:

```
$ sprof -c libdemo.so.1 $LD_PROFILE_OUTPUT/libdemo.so.1.profile
<UNKNOWN>                                x1                                1
```

x1	consumeCpu1	100
<UNKNOWN>	x2	1
x2	consumeCpu2	1000

**SEE ALSO****gprof(1), ldd(1), ld.so(8)****COLOPHON**

This page is part of release 5.02 of the Linux *man-pages* project. A description of the project, information about reporting bugs, and the latest version of this page, can be found at <https://www.kernel.org/doc/man-pages/>.