

NAME

`symlink`, `symlinkat` – make a new name for a file

SYNOPSIS

```
#include <unistd.h>
```

```
int symlink(const char *target, const char *linkpath);
```

```
#include <fcntl.h>      /* Definition of AT_* constants */
```

```
#include <unistd.h>
```

```
int symlinkat(const char *target, int newdirfd, const char *linkpath);
```

Feature Test Macro Requirements for glibc (see `feature_test_macros(7)`):

```
symlink():
```

```
_XOPEN_SOURCE >= 500 || _POSIX_C_SOURCE >= 200112L
|| /* Glibc versions <= 2.19: */ _BSD_SOURCE
```

```
symlinkat():
```

```
Since glibc 2.10:
```

```
_POSIX_C_SOURCE >= 200809L
```

```
Before glibc 2.10:
```

```
_ATFILE_SOURCE
```

DESCRIPTION

`symlink()` creates a symbolic link named *linkpath* which contains the string *target*.

Symbolic links are interpreted at run time as if the contents of the link had been substituted into the path being followed to find a file or directory.

Symbolic links may contain `..` path components, which (if used at the start of the link) refer to the parent directories of that in which the link resides.

A symbolic link (also known as a soft link) may point to an existing file or to a nonexistent one; the latter case is known as a dangling link.

The permissions of a symbolic link are irrelevant; the ownership is ignored when following the link, but is checked when removal or renaming of the link is requested and the link is in a directory with the sticky bit (`S_ISVTX`) set.

If *linkpath* exists, it will *not* be overwritten.

symlinkat()

The `symlinkat()` system call operates in exactly the same way as `symlink()`, except for the differences described here.

If the pathname given in *linkpath* is relative, then it is interpreted relative to the directory referred to by the file descriptor *newdirfd* (rather than relative to the current working directory of the calling process, as is done by `symlink()` for a relative pathname).

If *linkpath* is relative and *newdirfd* is the special value `AT_FDCWD`, then *linkpath* is interpreted relative to the current working directory of the calling process (like `symlink()`).

If *linkpath* is absolute, then *newdirfd* is ignored.

RETURN VALUE

On success, zero is returned. On error, `-1` is returned, and *errno* is set appropriately.

ERRORS**EACCES**

Write access to the directory containing *linkpath* is denied, or one of the directories in the path prefix of *linkpath* did not allow search permission. (See also `path_resolution(7)`.)

EDQUOT

The user's quota of resources on the filesystem has been exhausted. The resources could be inodes or disk blocks, depending on the filesystem implementation.

EEXIST

linkpath already exists.

EFAULT

target or *linkpath* points outside your accessible address space.

EIO An I/O error occurred.

ELOOP

Too many symbolic links were encountered in resolving *linkpath*.

ENAMETOOLONG

target or *linkpath* was too long.

ENOENT

A directory component in *linkpath* does not exist or is a dangling symbolic link, or *target* or *linkpath* is an empty string.

ENOMEM

Insufficient kernel memory was available.

ENOSPC

The device containing the file has no room for the new directory entry.

ENOTDIR

A component used as a directory in *linkpath* is not, in fact, a directory.

EPERM

The filesystem containing *linkpath* does not support the creation of symbolic links.

EROFS

linkpath is on a read-only filesystem.

The following additional errors can occur for **symlinkat()**:

EBADF

newdirfd is not a valid file descriptor.

ENOENT

linkpath is a relative pathname and *newdirfd* refers to a directory that has been deleted.

ENOTDIR

linkpath is relative and *newdirfd* is a file descriptor referring to a file other than a directory.

VERSIONS

symlinkat() was added to Linux in kernel 2.6.16; library support was added to glibc in version 2.4.

CONFORMING TO

symlink(): SVr4, 4.3BSD, POSIX.1-2001, POSIX.1-2008.

symlinkat(): POSIX.1-2008.

NOTES

No checking of *target* is done.

Deleting the name referred to by a symbolic link will actually delete the file (unless it also has other hard links). If this behavior is not desired, use **link(2)**.

Glibc notes

On older kernels where **symlinkat()** is unavailable, the glibc wrapper function falls back to the use of **symlink()**. When *linkpath* is a relative pathname, glibc constructs a pathname based on the symbolic link in */proc/self/fd* that corresponds to the *newdirfd* argument.

SEE ALSO

ln(1), namei(1), lchown(2), link(2), lstat(2), open(2), readlink(2), rename(2), unlink(2), path_resolution(7), symlink(7)

COLOPHON

This page is part of release 5.02 of the Linux *man-pages* project. A description of the project, information about reporting bugs, and the latest version of this page, can be found at <https://www.kernel.org/doc/man-pages/>.