

NAME

unshare – run program with some namespaces unshared from parent

SYNOPSIS

unshare [options] [*program* [*arguments*]]

DESCRIPTION

Unshares the indicated namespaces from the parent process and then executes the specified *program*. If *program* is not given, then “`${SHELL}`” is run (default: `/bin/sh`).

The namespaces can optionally be made persistent by bind mounting `/proc/pid/ns/type` files to a filesystem path and entered with **nsenter**(1) even after the *program* terminates (except PID namespaces where permanently running init process is required). Once a persistent namespace is no longer needed, it can be unperisted with **umount**(8). See the **EXAMPLES** section for more details.

The namespaces to be unshared are indicated via options. Unshareable namespaces are:

mount namespace

Mounting and unmounting filesystems will not affect the rest of the system, except for filesystems which are explicitly marked as shared (with **mount --make-shared**; see `/proc/self/mountinfo` or **findmnt -o+PROPAGATION** for the **shared** flags). For further details, see **mount_namespaces**(7) and the discussion of the **CLONE_NEWNS** flag in **clone**(2).

unshare since util-linux version 2.27 automatically sets propagation to **private** in a new mount namespace to make sure that the new namespace is really unshared. It’s possible to disable this feature with option **--propagation unchanged**. Note that **private** is the kernel default.

UTS namespace

Setting hostname or domainname will not affect the rest of the system. For further details, see **namespaces**(7) and the discussion of the **CLONE_NEWUTS** flag in **clone**(2).

IPC namespace

The process will have an independent namespace for POSIX message queues as well as System V message queues, semaphore sets and shared memory segments. For further details, see **namespaces**(7) and the discussion of the **CLONE_NEWIPC** flag in **clone**(2).

network namespace

The process will have independent IPv4 and IPv6 stacks, IP routing tables, firewall rules, the `/proc/net` and `/sys/class/net` directory trees, sockets, etc. For further details, see **namespaces**(7) and the discussion of the **CLONE_NEWNET** flag in **clone**(2).

PID namespace

Children will have a distinct set of PID-to-process mappings from their parent. For further details, see **pid_namespaces**(7) and the discussion of the **CLONE_NEWPID** flag in **clone**(2).

cgroup namespace

The process will have a virtualized view of `/proc/self/cgroup`, and new cgroup mounts will be rooted at the namespace cgroup root. For further details, see **cgroup_namespaces**(7) and the discussion of the **CLONE_NEWCGROUP** flag in **clone**(2).

user namespace

The process will have a distinct set of UIDs, GIDs and capabilities. For further details, see **user_namespaces**(7) and the discussion of the **CLONE_NEWUSER** flag in **clone**(2).

OPTIONS

-i, --ipc[=*file*]

Unshare the IPC namespace. If *file* is specified, then a persistent namespace is created by a bind mount.

-m, --mount[=*file*]

Unshare the mount namespace. If *file* is specified, then a persistent namespace is created by a bind mount. Note that *file* has to be located on a filesystem with the propagation flag set to **private**.

Use the command **findmnt -o+PROPAGATION** when not sure about the current setting. See also the examples below.

- n, --net[=*file*]**
Unshare the network namespace. If *file* is specified, then a persistent namespace is created by a bind mount.
- p, --pid[=*file*]**
Unshare the PID namespace. If *file* is specified then persistent namespace is created by a bind mount. See also the **--fork** and **--mount-proc** options.
- u, --uts[=*file*]**
Unshare the UTS namespace. If *file* is specified, then a persistent namespace is created by a bind mount.
- U, --user[=*file*]**
Unshare the user namespace. If *file* is specified, then a persistent namespace is created by a bind mount.
- C, --cgroup[=*file*]**
Unshare the cgroup namespace. If *file* is specified then persistent namespace is created by bind mount.
- f, --fork**
Fork the specified *program* as a child process of **unshare** rather than running it directly. This is useful when creating a new PID namespace.
- kill-child[=*signame*]**
When **unshare** terminates, have *signame* be sent to the forked child process. Combined with **--pid** this allows for an easy and reliable killing of the entire process tree below **unshare**. If not given, *signame* defaults to **SIGKILL**. This option implies **--fork**.
- mount-proc[=*mountpoint*]**
Just before running the program, mount the proc filesystem at *mountpoint* (default is /proc). This is useful when creating a new PID namespace. It also implies creating a new mount namespace since the /proc mount would otherwise mess up existing programs on the system. The new proc filesystem is explicitly mounted as private (with MS_PRIVATE|MS_REC).
- r, --map-root-user**
Run the program only after the current effective user and group IDs have been mapped to the superuser UID and GID in the newly created user namespace. This makes it possible to conveniently gain capabilities needed to manage various aspects of the newly created namespaces (such as configuring interfaces in the network namespace or mounting filesystems in the mount namespace) even when run unprivileged. As a mere convenience feature, it does not support more sophisticated use cases, such as mapping multiple ranges of UIDs and GIDs. This option implies **--setgroups=deny**.
- propagation private|shared|slave|unchanged**
Recursively set the mount propagation flag in the new mount namespace. The default is to set the propagation to *private*. It is possible to disable this feature with the argument **unchanged**. The option is silently ignored when the mount namespace (**--mount**) is not requested.
- setgroups allow|deny**
Allow or deny the **setgroups(2)** system call in a user namespace.

To be able to call **setgroups(2)**, the calling process must at least have CAP_SETGID. But since Linux 3.19 a further restriction applies: the kernel gives permission to call **setgroups(2)** only after the GID map (**/proc/pid/gid_map**) has been set. The GID map is writable by root when **setgroups(2)** is enabled (i.e. **allow**, the default), and the GID map becomes writable by unprivileged processes when **setgroups(2)** is permanently disabled (with **deny**).

- R, --root=*dir***
run the command with root directory set to *dir*.
- w, --wd=*dir***
change working directory to *dir*.
- S, --setuid *uid***
Set the user ID which will be used in the entered namespace.
- G, --setgid *gid***
Set the group ID which will be used in the entered namespace and drop supplementary groups.
- V, --version**
Display version information and exit.
- h, --help**
Display help text and exit.

NOTES

The proc and sysfs filesystems mounting as root in a user namespace have to be restricted so that a less privileged user can not get more access to sensitive files that a more privileged user made unavailable. In short the rule for proc and sysfs is as close to a bind mount as possible.

EXAMPLES

```
# unshare --fork --pid --mount-proc readlink /proc/self
```

```
1
```

Establish a PID namespace, ensure we're PID 1 in it against a newly mounted procfs instance.

```
$ unshare --map-root-user --user sh -c whoami
```

```
root
```

Establish a user namespace as an unprivileged user with a root user within it.

```
# touch /root/uts-ns
```

```
# unshare --uts=/root/uts-ns hostname FOO
```

```
# nsenter --uts=/root/uts-ns hostname
```

```
FOO
```

```
# umount /root/uts-ns
```

Establish a persistent UTS namespace, and modify the hostname. The namespace is then entered with **nsenter**. The namespace is destroyed by unmounting the bind reference.

```
# mount --bind /root/namespaces /root/namespaces
```

```
# mount --make-private /root/namespaces
```

```
# touch /root/namespaces/mnt
```

```
# unshare --mount=/root/namespaces/mnt
```

Establish a persistent mount namespace referenced by the bind mount /root/namespaces/mnt. This example shows a portable solution, because it makes sure that the bind mount is created on a shared filesystem.

```
# unshare -pf --kill-child -- bash -c (sleep 999 &) && sleep 1000 &
```

```
# pid=$!
```

```
# kill $pid
```

Reliable killing of subprocesses of the *program*. When **unshare** gets killed, everything below it gets killed as well. Without it, the children of *program* would have orphaned and been re-parented to PID 1.

SEE ALSO

clone(2), **unshare(2)**, **namespaces(7)**, **mount(8)**

AUTHORS

Mikhail Gusarov <dottedmag@dottedmag.net>

Karel Zak <kzak@redhat.com>

AVAILABILITY

The `unshare` command is part of the `util-linux` package and is available from <https://www.kernel.org/pub/linux/utils/util-linux/>.